

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
import re
from sklearn.datasets import load_digits
```

```
In [2]: a=pd.read_csv(r"C:\Users\user\Downloads\C10_air\csvs_per_year\csvs_per_year\ma
```

Out[2]:

	date	BEN	CO	EBE	MXV	NMHC	NO_2	NOx	OXY	O_3	
0	2003-03-01 01:00:00	NaN	1.72	NaN	NaN	NaN	73.900002	316.299988	NaN	10.550000	55.2
1	2003-03-01 01:00:00	NaN	1.45	NaN	NaN	0.26	72.110001	250.000000	0.73	6.720000	52.3
2	2003-03-01 01:00:00	NaN	1.57	NaN	NaN	NaN	80.559998	224.199997	NaN	21.049999	63.2
3	2003-03-01 01:00:00	NaN	2.45	NaN	NaN	NaN	78.370003	450.399994	NaN	4.220000	67.8
4	2003-03-01 01:00:00	NaN	3.26	NaN	NaN	NaN	96.250000	479.100006	NaN	8.460000	95.7
...	
243979	2003-10-01 00:00:00	0.20	0.16	2.01	3.17	0.02	31.799999	32.299999	1.68	34.049999	7.3
243980	2003-10-01 00:00:00	0.32	0.08	0.36	0.72	NaN	10.450000	14.760000	1.00	34.610001	7.4
243981	2003-10-01 00:00:00	NaN	NaN	NaN	NaN	0.07	34.639999	50.810001	NaN	32.160000	16.8
243982	2003-10-01 00:00:00	NaN	NaN	NaN	NaN	0.07	32.580002	41.020000	NaN	NaN	13.5
243983	2003-10-01 00:00:00	1.00	0.29	2.15	6.41	0.07	37.150002	56.849998	2.28	21.480000	12.3

243984 rows × 16 columns

In [3]:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 243984 entries, 0 to 243983
Data columns (total 16 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        243984 non-null  object
1   BEN         69745 non-null   float64
2   CO          225340 non-null  float64
3   EBE         61244 non-null   float64
4   MXY         42045 non-null   float64
5   NMHC        111951 non-null   float64
6   NO_2        242625 non-null   float64
7   NOx         242629 non-null   float64
8   OXY         42072 non-null   float64
9   O_3         234131 non-null   float64
10  PM10        240896 non-null   float64
11  PXY         42063 non-null   float64
12  SO_2        242729 non-null   float64
13  TCH         111991 non-null   float64
14  TOL         69439 non-null   float64
15  station     243984 non-null   int64
dtypes: float64(14), int64(1), object(1)
memory usage: 29.8+ MB
```

```
In [4]: b=a.fillna(value=126)
```

```
Out[4]:
```

	date	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	
0	2003-03-01 01:00:00	126.00	1.72	126.00	126.00	126.00	73.900002	316.299988	126.00	10.51
1	2003-03-01 01:00:00	126.00	1.45	126.00	126.00	0.26	72.110001	250.000000	0.73	6.72
2	2003-03-01 01:00:00	126.00	1.57	126.00	126.00	126.00	80.559998	224.199997	126.00	21.04
3	2003-03-01 01:00:00	126.00	2.45	126.00	126.00	126.00	78.370003	450.399994	126.00	4.22
4	2003-03-01 01:00:00	126.00	3.26	126.00	126.00	126.00	96.250000	479.100006	126.00	8.46
...
243979	2003-10-01 00:00:00	0.20	0.16	2.01	3.17	0.02	31.799999	32.299999	1.68	34.04
243980	2003-10-01 00:00:00	0.32	0.08	0.36	0.72	126.00	10.450000	14.760000	1.00	34.61
243981	2003-10-01 00:00:00	126.00	126.00	126.00	126.00	0.07	34.639999	50.810001	126.00	32.16
243982	2003-10-01 00:00:00	126.00	126.00	126.00	126.00	0.07	32.580002	41.020000	126.00	126.00
243983	2003-10-01 00:00:00	1.00	0.29	2.15	6.41	0.07	37.150002	56.849998	2.28	21.48

243984 rows × 16 columns

```
In [5]:
```

```
Out[5]: Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',  
'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'],  
             dtype='object')
```

```
In [6]: c=b.head(11)
```

Out[6]:

	date	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3
0	2003-03-01 01:00:00	126.00	1.72	126.00	126.00	126.00	73.900002	316.299988	126.00	10.550000
1	2003-03-01 01:00:00	126.00	1.45	126.00	126.00	0.26	72.110001	250.000000	0.73	6.720000
2	2003-03-01 01:00:00	126.00	1.57	126.00	126.00	126.00	80.559998	224.199997	126.00	21.049999
3	2003-03-01 01:00:00	126.00	2.45	126.00	126.00	126.00	78.370003	450.399994	126.00	4.220000
4	2003-03-01 01:00:00	126.00	3.26	126.00	126.00	126.00	96.250000	479.100006	126.00	8.460000
5	2003-03-01 01:00:00	8.41	1.94	9.83	21.49	0.45	90.300003	384.899994	9.48	9.950000
6	2003-03-01 01:00:00	126.00	1.38	126.00	126.00	0.29	89.580002	230.000000	126.00	7.200000
7	2003-03-01 01:00:00	126.00	1.58	126.00	126.00	0.30	93.639999	334.600006	126.00	4.190000
8	2003-03-01 01:00:00	126.00	126.00	126.00	126.00	126.00	126.000000	126.000000	126.00	126.000000
9	2003-03-01 01:00:00	126.00	1.92	126.00	126.00	126.00	71.839996	181.399994	126.00	5.330000
10	2003-03-01 01:00:00	126.00	1.33	126.00	126.00	0.31	87.919998	273.399994	126.00	4.250000


```
In [7]: d=c[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3']]
```

Out[7]:

	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3
0	126.00	1.72	126.00	126.00	126.00	73.900002	316.299988	126.00	10.550000
1	126.00	1.45	126.00	126.00	0.26	72.110001	250.000000	0.73	6.720000
2	126.00	1.57	126.00	126.00	126.00	80.559998	224.199997	126.00	21.049999
3	126.00	2.45	126.00	126.00	126.00	78.370003	450.399994	126.00	4.220000
4	126.00	3.26	126.00	126.00	126.00	96.250000	479.100006	126.00	8.460000
5	8.41	1.94	9.83	21.49	0.45	90.300003	384.899994	9.48	9.950000
6	126.00	1.38	126.00	126.00	0.29	89.580002	230.000000	126.00	7.200000
7	126.00	1.58	126.00	126.00	0.30	93.639999	334.600006	126.00	4.190000
8	126.00	126.00	126.00	126.00	126.00	126.000000	126.000000	126.00	126.000000
9	126.00	1.92	126.00	126.00	126.00	71.839996	181.399994	126.00	5.330000
10	126.00	1.33	126.00	126.00	0.31	87.919998	273.399994	126.00	4.250000


In [8]:

Out[8]: <AxesSubplot:>



In [9]:

Out[9]: <seaborn.axisgrid.PairGrid at 0x1fdea615250>

In [10]: `x=d[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY']]`In [11]: `from sklearn.model_selection import train_test_split`In [12]: `from sklearn.linear_model import LinearRegression`
`lr=LinearRegression()`Out[12]: `LinearRegression()`

In [13]:

1.0231815394945443e-12


```
In [14]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
```

```
Out[14]:
```

	Co-efficient
BEN	-1.695639e-16
CO	1.653345e-14
EBE	-3.975580e-16
MXY	4.284913e-16
NMHC	-4.011195e-15
NO_2	-2.801221e-14
NOx	1.000000e+00
OXY	2.871689e-15

```
In [15]: prediction=lr.predict(x_test)
```

```
Out[15]: <matplotlib.collections.PathCollection at 0x1fdef63b220>
```



```
In [16]:
```

```
1.0
```

```
In [17]:
```

```
In [18]: rr=Ridge(alpha=10)
```

```
Out[18]: Ridge(alpha=10)
```

```
In [19]:
```

```
Out[19]: 0.9998869821152107
```

```
In [20]: la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

```
Out[20]: Lasso(alpha=10)
```

In [21]:

Out[21]: 0.9999989799894025

In [22]: a1=b.head(5000)

Out[22]:

	date	BEN	CO	EBE	MXV	NMHC	NO_2	NOx	OXY	O_3
0	2003-03-01 01:00:00	126.00	1.72	126.00	126.0	126.00	73.900002	316.299988	126.00	10.550000
1	2003-03-01 01:00:00	126.00	1.45	126.00	126.0	0.26	72.110001	250.000000	0.73	6.720000
2	2003-03-01 01:00:00	126.00	1.57	126.00	126.0	126.00	80.559998	224.199997	126.00	21.049999
3	2003-03-01 01:00:00	126.00	2.45	126.00	126.0	126.00	78.370003	450.399994	126.00	4.220000
4	2003-03-01 01:00:00	126.00	3.26	126.00	126.0	126.00	96.250000	479.100006	126.00	8.460000
...
4995	2003-03-08 11:00:00	126.00	0.46	126.00	126.0	126.00	27.469999	36.110001	126.00	43.799999
4996	2003-03-08 11:00:00	126.00	0.69	126.00	126.0	0.07	39.020000	56.230000	126.00	39.720001
4997	2003-03-08 11:00:00	126.00	1.74	126.00	126.0	126.00	65.790001	148.500000	126.00	23.760000
4998	2003-03-08 11:00:00	1.85	0.56	1.62	126.0	0.13	52.610001	96.300003	126.00	28.770000
4999	2003-03-08 11:00:00	126.00	0.84	126.00	126.0	126.00	33.310001	41.720001	126.00	38.459999

5000 rows × 16 columns

In [23]: e=a1[['BEN', 'CO', 'EBE', 'MXV', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',

In [24]: f=e.iloc[:,0:14]

In [25]:

In [26]: logr=LogisticRegression(max_iter=10000)

Out[26]: LogisticRegression(max_iter=10000)

In [27]: from sklearn.model_selection import train_test_split

In [28]:


```
In [29]: prediction=logr.predict(i)
```

```
[28079003]
```

```
In [30]:
```

```
Out[30]: array([28079001, 28079003, 28079004, 28079006, 28079007, 28079008,
                28079009, 28079011, 28079012, 28079014, 28079015, 28079016,
                28079017, 28079018, 28079019, 28079021, 28079022, 28079023,
                28079024, 28079025, 28079026, 28079027, 28079035, 28079036,
                28079038, 28079039, 28079040, 28079099], dtype=int64)
```

```
In [31]:
```

```
Out[31]: 9.896956293105757e-13
```

```
In [32]:
```

```
Out[32]: 0.9999999999683131
```

```
In [33]:
```

```
Out[33]: 0.594
```

```
In [34]: from sklearn.linear_model import ElasticNet
          en=ElasticNet()
```

```
Out[34]: ElasticNet()
```

```
In [35]:
```

```
[ -0.          -0.          -0.          -0.          -0.          -0.
  0.99990495 -0.          ]
```

```
In [36]:
```

```
0.029388644480661696
```

```
In [37]: prediction=en.predict(x_test)
```

```
0.9999999898008636
```

```
In [38]: from sklearn.ensemble import RandomForestClassifier
          rfc=RandomForestClassifier()
```

```
Out[38]: RandomForestClassifier()
```

```
In [39]: parameters={'max_depth':[1,2,3,4,5],
                    'min_samples_leaf':[5,10,15,20,25],
                    'n_estimators':[10,20,30,40,50]}
```

```
In [40]: from sklearn.model_selection import GridSearchCV
grid_search=GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="acc
```

```
Out[40]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
                    param_grid={'max_depth': [1, 2, 3, 4, 5],
                                'min_samples_leaf': [5, 10, 15, 20, 25],
                                'n_estimators': [10, 20, 30, 40, 50]},
                    scoring='accuracy')
```

```
In [41]:
```

```
Out[41]: 0.5691428571428572
```

```
In [42]:
```

```
In [43]: from sklearn.tree import plot_tree
plt.figure(figsize=(80,50))
```

```
Out[43]: [Text(1900.2162162162163, 2491.5, 'X[0] <= -0.355\ngini = 0.964\nsamples = 22
24\nvalue = [126, 136, 138, 119, 113, 99, 117, 120, 124, 137\n106, 131, 127,
135, 114, 127, 146, 129, 108, 109\n127, 103, 129, 142, 147, 131, 123, 137]'),
Text(573.081081081081, 2038.5, 'X[11] <= -0.864\ngini = 0.872\nsamples = 59
1\nvalue = [0, 0, 0, 119, 0, 0, 0, 0, 0, 0, 106, 0, 0, 0\n0, 0, 137, 129, 10
8, 109, 0, 0, 75, 0, 0, 0, 0\n136]'),
Text(241.2972972972973, 1585.5, 'X[10] <= -2.287\ngini = 0.157\nsamples = 4
2\nvalue = [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0\n0, 0, 0, 6, 64, 0, 0,
0, 0, 0, 0, 0, 0]'),
Text(120.64864864864865, 1132.5, 'gini = 0.0\nsamples = 37\nvalue = [0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0\n0, 0, 0, 0, 62, 0, 0, 0, 0, 0, 0, 0,
0]'),
Text(361.94594594594594, 1132.5, 'gini = 0.375\nsamples = 5\nvalue = [0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0\n0, 0, 0, 6, 2, 0, 0, 0, 0, 0, 0, 0,
0]'),
Text(904.8648648648648, 1585.5, 'X[7] <= -0.661\ngini = 0.865\nsamples = 54
9\nvalue = [0, 0, 0, 119, 0, 0, 0, 0, 0, 0, 0, 0, 106, 0, 0, 0\n0, 0, 137, 123, 44,
109, 0, 0, 75, 0, 0, 0, 0\n136]'),
Text(603.2432432432432, 1132.5, 'X[4] <= -0.069\ngini = 0.777\nsamples = 31
0\nvalue = [0, 0, 0, 119, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0\n0, 0, 0, 0, 11, 100
```

**From this observation I had observe that the
ELASTICNET highest accuracy of
0.9999999898008636**

```
In [ ]:
```