

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
import re
from sklearn.datasets import load_digits
```

```
In [2]: a=pd.read_csv(r"C:\Users\user\Downloads\C10_air\csvs_per_year\csvs_per_year\ma
```

Out[2]:

	date	BEN	CO	EBE	MXV	NMHC	NO_2	NOx	OXY	O_3	
0	2004-08-01 01:00:00	NaN	0.66	NaN	NaN	NaN	89.550003	118.900002	NaN	40.020000	39.
1	2004-08-01 01:00:00	2.66	0.54	2.99	6.08	0.18	51.799999	53.860001	3.28	51.689999	22.
2	2004-08-01 01:00:00	NaN	1.02	NaN	NaN	NaN	93.389999	138.600006	NaN	20.860001	49.
3	2004-08-01 01:00:00	NaN	0.53	NaN	NaN	NaN	87.290001	105.000000	NaN	36.730000	31.
4	2004-08-01 01:00:00	NaN	0.17	NaN	NaN	NaN	34.910000	35.349998	NaN	86.269997	54.
...	...	...	...	...	...	...	...	...	...	...	...
245491	2004-06-01 00:00:00	0.75	0.21	0.85	1.55	0.07	59.580002	64.389999	0.66	33.029999	30.
245492	2004-06-01 00:00:00	2.49	0.75	2.44	4.57	NaN	97.139999	146.899994	2.34	7.740000	37.
245493	2004-06-01 00:00:00	NaN	NaN	NaN	NaN	0.13	102.699997	132.600006	NaN	17.809999	22.
245494	2004-06-01 00:00:00	NaN	NaN	NaN	NaN	0.09	82.599998	102.599998	NaN	NaN	45.
245495	2004-06-01 00:00:00	3.01	0.67	2.78	5.12	0.20	92.550003	141.000000	2.60	11.460000	24.

245496 rows × 17 columns

In [3]:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 245496 entries, 0 to 245495
Data columns (total 17 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        245496 non-null object
1   BEN         65158 non-null float64
2   CO          226043 non-null float64
3   EBE         56781 non-null float64
4   MXY         39867 non-null float64
5   NMHC        107630 non-null float64
6   NO_2        243280 non-null float64
7   NOx         243283 non-null float64
8   OXY         39882 non-null float64
9   O_3         233811 non-null float64
10  PM10        234655 non-null float64
11  PM25        58145 non-null float64
12  PXY         39891 non-null float64
13  SO_2        243402 non-null float64
14  TCH         107650 non-null float64
15  TOL         64914 non-null float64
16  station     245496 non-null int64
dtypes: float64(15), int64(1), object(1)
memory usage: 31.8+ MB
```

```
In [4]: b=a.fillna(value=106)
```

```
Out[4]:
```

	date	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	
0	2004-08-01 01:00:00	106.00	0.66	106.00	106.00	106.00	89.550003	118.900002	106.00	40.0
1	2004-08-01 01:00:00	2.66	0.54	2.99	6.08	0.18	51.799999	53.860001	3.28	51.6
2	2004-08-01 01:00:00	106.00	1.02	106.00	106.00	106.00	93.389999	138.600006	106.00	20.8
3	2004-08-01 01:00:00	106.00	0.53	106.00	106.00	106.00	87.290001	105.000000	106.00	36.7
4	2004-08-01 01:00:00	106.00	0.17	106.00	106.00	106.00	34.910000	35.349998	106.00	86.2
...	...	...	...	...	...	...	...	...	...	...
245491	2004-06-01 00:00:00	0.75	0.21	0.85	1.55	0.07	59.580002	64.389999	0.66	33.0
245492	2004-06-01 00:00:00	2.49	0.75	2.44	4.57	106.00	97.139999	146.899994	2.34	7.7
245493	2004-06-01 00:00:00	106.00	106.00	106.00	106.00	0.13	102.699997	132.600006	106.00	17.8
245494	2004-06-01 00:00:00	106.00	106.00	106.00	106.00	0.09	82.599998	102.599998	106.00	106.0
245495	2004-06-01 00:00:00	3.01	0.67	2.78	5.12	0.20	92.550003	141.000000	2.60	11.4

245496 rows × 17 columns

```
In [5]:
```

```
Out[5]: Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',  
'PM10', 'PM25', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'],  
              dtype='object')
```

```
In [6]: c=b.head(11)
```

```
Out[6]:
```

	date	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3
0	2004-08-01 01:00:00	106.00	0.66	106.00	106.00	106.00	89.550003	118.900002	106.00	40.020000
1	2004-08-01 01:00:00	2.66	0.54	2.99	6.08	0.18	51.799999	53.860001	3.28	51.689999
2	2004-08-01 01:00:00	106.00	1.02	106.00	106.00	106.00	93.389999	138.600006	106.00	20.860001
3	2004-08-01 01:00:00	106.00	0.53	106.00	106.00	106.00	87.290001	105.000000	106.00	36.730000
4	2004-08-01 01:00:00	106.00	0.17	106.00	106.00	106.00	34.910000	35.349998	106.00	86.269997
5	2004-08-01 01:00:00	3.24	0.63	5.55	9.72	0.06	103.800003	144.800003	5.04	32.480000
6	2004-08-01 01:00:00	106.00	0.43	106.00	106.00	0.17	54.270000	64.279999	106.00	66.589996
7	2004-08-01 01:00:00	1.41	0.47	2.35	106.00	0.02	71.730003	87.519997	106.00	53.270000
8	2004-08-01 01:00:00	106.00	1.28	106.00	106.00	106.00	147.699997	202.500000	106.00	10.280000
9	2004-08-01 01:00:00	106.00	0.43	106.00	106.00	0.27	54.290001	68.099998	106.00	66.709999
10	2004-08-01 01:00:00	106.00	0.60	106.00	106.00	106.00	73.410004	87.059998	106.00	40.480000


```
In [7]: d=c[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3']]
```

```
Out[7]:
```

	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3
0	106.00	0.66	106.00	106.00	106.00	89.550003	118.900002	106.00	40.020000
1	2.66	0.54	2.99	6.08	0.18	51.799999	53.860001	3.28	51.689999
2	106.00	1.02	106.00	106.00	106.00	93.389999	138.600006	106.00	20.860001
3	106.00	0.53	106.00	106.00	106.00	87.290001	105.000000	106.00	36.730000
4	106.00	0.17	106.00	106.00	106.00	34.910000	35.349998	106.00	86.269997
5	3.24	0.63	5.55	9.72	0.06	103.800003	144.800003	5.04	32.480000
6	106.00	0.43	106.00	106.00	0.17	54.270000	64.279999	106.00	66.589996
7	1.41	0.47	2.35	106.00	0.02	71.730003	87.519997	106.00	53.270000
8	106.00	1.28	106.00	106.00	106.00	147.699997	202.500000	106.00	10.280000
9	106.00	0.43	106.00	106.00	0.27	54.290001	68.099998	106.00	66.709999
10	106.00	0.60	106.00	106.00	106.00	73.410004	87.059998	106.00	40.480000


In [8]:

Out[8]: &lt;AxesSubplot:&gt;



In [9]:

Out[9]: &lt;seaborn.axisgrid.PairGrid at 0x2192516a160&gt;

In [10]: `x=d[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY']]`In [11]: `from sklearn.model_selection import train_test_split`In [12]: `from sklearn.linear_model import LinearRegression`  
`lr=LinearRegression()`

Out[12]: LinearRegression()

In [13]:

-0.0008645494699663914


```
In [14]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
```

```
Out[14]:
```

	Co-efficient
<b>BEN</b>	0.000833
<b>CO</b>	-0.012756
<b>EBE</b>	-0.000813
<b>MXY</b>	-0.000835
<b>NMHC</b>	-0.000017
<b>NO_2</b>	0.000374
<b>NOx</b>	0.999834
<b>OXY</b>	0.000790

```
In [15]: prediction=lr.predict(x_test)
```

```
Out[15]: <matplotlib.collections.PathCollection at 0x2192dbc92b0>
```



```
In [16]:
```

```
0.9999999916628053
```

```
In [17]:
```

```
In [18]: rr=Ridge(alpha=10)
```

```
Out[18]: Ridge(alpha=10)
```

```
In [19]:
```

```
Out[19]: 0.9993279789663293
```

```
In [20]: la=Lasso(alpha=10)
```

```
la.fit(x_train,y_train)
```

```
Out[20]: Lasso(alpha=10)
```

In [21]:

Out[21]: 0.9999734689666656

In [22]: a1=b.head(4000)

Out[22]:

	date	BEN	CO	EBE	MXV	NMHC	NO_2	NOx	OXY	O_3
0	2004-08-01 01:00:00	106.00	0.66	106.00	106.00	106.00	89.550003	118.900002	106.00	40.020000
1	2004-08-01 01:00:00	2.66	0.54	2.99	6.08	0.18	51.799999	53.860001	3.28	51.689999
2	2004-08-01 01:00:00	106.00	1.02	106.00	106.00	106.00	93.389999	138.600006	106.00	20.860000
3	2004-08-01 01:00:00	106.00	0.53	106.00	106.00	106.00	87.290001	105.000000	106.00	36.730000
4	2004-08-01 01:00:00	106.00	0.17	106.00	106.00	106.00	34.910000	35.349998	106.00	86.269997
...	...	...	...	...	...	...	...	...	...	...
3995	2004-08-07 00:00:00	106.00	0.75	106.00	106.00	106.00	114.900002	189.300003	106.00	12.330000
3996	2004-08-07 00:00:00	106.00	1.68	106.00	106.00	106.00	143.199997	321.000000	106.00	8.710000
3997	2004-08-07 00:00:00	106.00	0.65	106.00	106.00	0.65	130.199997	156.500000	106.00	7.600000
3998	2004-08-07 00:00:00	106.00	0.82	106.00	106.00	106.00	125.400002	207.199997	106.00	9.000000
3999	2004-08-07 00:00:00	106.00	0.74	106.00	106.00	106.00	104.800003	121.599998	106.00	18.430000

4000 rows × 17 columns

In [23]: e=a1[['BEN', 'CO', 'EBE', 'MXV', 'NMHC', 'NO\_2', 'NOx', 'OXY', 'O\_3',

In [24]: f=e.iloc[:,0:14]

In [25]:

In [26]: logr=LogisticRegression(max\_iter=10000)

Out[26]: LogisticRegression(max\_iter=10000)

In [27]: from sklearn.model\_selection import train\_test\_split

In [28]:



```
In [29]: prediction=logr.predict(i)
```

```
[28079004]
```

```
In [30]:
```

```
Out[30]: array([28079001, 28079003, 28079004, 28079006, 28079007, 28079008,
                28079009, 28079011, 28079012, 28079014, 28079015, 28079016,
                28079017, 28079018, 28079019, 28079021, 28079022, 28079023,
                28079024, 28079025, 28079026, 28079027, 28079035, 28079036,
                28079038, 28079039, 28079040, 28079099], dtype=int64)
```

```
In [31]:
```

```
Out[31]: 2.793068192100045e-39
```

```
In [32]:
```

```
Out[32]: 5.075635610824596e-186
```

```
In [33]:
```

```
Out[33]: 0.6116666666666667
```

```
In [34]: from sklearn.linear_model import ElasticNet
          en=ElasticNet()
```

```
Out[34]: ElasticNet()
```

```
In [35]:
```

```
[ 0.00000000e+00 -0.00000000e+00  4.45095914e-04 -0.00000000e+00
 -7.95362564e-04  3.16610023e-02  9.79393568e-01 -0.00000000e+00]
```

```
In [36]:
```

```
-0.4250071227973251
```

```
In [37]: prediction=en.predict(x_test)
```

```
0.9999745139676599
```

```
In [38]: from sklearn.ensemble import RandomForestClassifier
          rfc=RandomForestClassifier()
```

```
Out[38]: RandomForestClassifier()
```

```
In [39]: parameters={'max_depth':[1,2,3,4,5],
                    'min_samples_leaf':[5,10,15,20,25],
                    'n_estimators':[10,20,30,40,50]}
```

```
In [40]: from sklearn.model_selection import GridSearchCV
grid_search=GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="acc
```

```
Out[40]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
                    param_grid={'max_depth': [1, 2, 3, 4, 5],
                                'min_samples_leaf': [5, 10, 15, 20, 25],
                                'n_estimators': [10, 20, 30, 40, 50]},
                    scoring='accuracy')
```

```
In [41]:
```

```
Out[41]: 0.6492857142857142
```

```
In [42]:
```

```
In [43]: from sklearn.tree import plot_tree
plt.figure(figsize=(80,50))
```

```
Out[43]: [Text(2028.1153846153845, 2491.5, 'X[0] <= -0.435\ngini = 0.964\nsamples = 17
64\nvalue = [117, 102, 127, 102, 115, 128, 102, 104, 77, 90, 87\n84, 80, 87,
95, 88, 107, 92, 94, 82, 115, 96\n101, 109, 109, 100, 101, 109]'),
Text(965.7692307692307, 2038.5, 'X[3] <= -0.699\ngini = 0.873\nsamples = 49
9\nvalue = [0, 0, 0, 101, 0, 128, 0, 0, 0, 0, 87, 0, 0\n0, 0, 0, 102, 0, 94,
80, 0, 0, 101, 0, 0, 0\n0, 109]'),
Text(472.15384615384613, 1585.5, 'X[3] <= -2.184\ngini = 0.798\nsamples = 30
9\nvalue = [0, 0, 0, 101, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0\n0, 0, 0, 0, 94, 80,
0, 0, 101, 0, 0, 0, 0\n109]'),
Text(257.53846153846155, 1132.5, 'X[12] <= -0.07\ngini = 0.613\nsamples = 10
7\nvalue = [0, 0, 0, 12, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0\n0, 0, 0, 0, 80, 58, 0,
0, 0, 0, 0, 0, 11]'),
Text(171.69230769230768, 679.5, 'X[0] <= -1.597\ngini = 0.372\nsamples = 71\
nvalue = [0, 0, 0, 12, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0\n0, 0, 0, 0, 80, 0, 0, 0,
0, 0, 0, 0, 11]'),
Text(85.84615384615384, 226.5, 'gini = 0.185\nsamples = 61\nvalue = [0, 0,
0, 7, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0\n0, 0, 0, 0, 80, 0, 0, 0, 0, 0, 0, 0, 0,
2]'),
Text(257.53846153846155, 226.5, 'gini = 0.459\nsamples = 10\nvalue = [0, 0,
```

**From this observation I had observe that the LINEAR  
REFRESSION highest accuracy of  
0.9999999916628053**

```
In [ ]:
```