

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
import re
from sklearn.datasets import load_digits
```

```
In [2]: a=pd.read_csv(r"C:\Users\user\Downloads\C10_air\csvs_per_year\csvs_per_year\ma
```

Out[2]:

	date	BEN	CO	EBE	MXV	NMHC	NO_2	NOx	OXY	O_3	
0	2008-06-01 01:00:00	NaN	0.47	NaN	NaN	NaN	83.089996	120.699997	NaN	16.990000	16.
1	2008-06-01 01:00:00	NaN	0.59	NaN	NaN	NaN	94.820000	130.399994	NaN	17.469999	19.
2	2008-06-01 01:00:00	NaN	0.55	NaN	NaN	NaN	75.919998	104.599998	NaN	13.470000	20.
3	2008-06-01 01:00:00	NaN	0.36	NaN	NaN	NaN	61.029999	66.559998	NaN	23.110001	10.
4	2008-06-01 01:00:00	1.68	0.80	1.70	3.01	0.30	105.199997	214.899994	1.61	12.120000	37.
...
226387	2008-11-01 00:00:00	0.48	0.30	0.57	1.00	0.31	13.050000	14.160000	0.91	57.400002	5.
226388	2008-11-01 00:00:00	NaN	0.30	NaN	NaN	NaN	41.880001	48.500000	NaN	35.830002	15.
226389	2008-11-01 00:00:00	0.25	NaN	0.56	NaN	0.11	83.610001	102.199997	NaN	14.130000	17.
226390	2008-11-01 00:00:00	0.54	NaN	2.70	NaN	0.18	70.639999	81.860001	NaN	NaN	11.
226391	2008-11-01 00:00:00	0.75	0.36	1.20	2.75	0.16	58.240002	74.239998	1.64	31.910000	12.

226392 rows × 17 columns

In [3]:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 226392 entries, 0 to 226391
Data columns (total 17 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        226392 non-null object
1   BEN         67047 non-null float64
2   CO          208109 non-null float64
3   EBE         67044 non-null float64
4   MXY         25867 non-null float64
5   NMHC        85079 non-null float64
6   NO_2        225315 non-null float64
7   NOx         225311 non-null float64
8   OXY         25878 non-null float64
9   O_3         215716 non-null float64
10  PM10        220179 non-null float64
11  PM25        67833 non-null float64
12  PXY         25877 non-null float64
13  SO_2        225405 non-null float64
14  TCH         85107 non-null float64
15  TOL         66940 non-null float64
16  station     226392 non-null int64
dtypes: float64(15), int64(1), object(1)
memory usage: 29.4+ MB
```

```
In [4]: b=a.fillna(value=119)
```

```
Out[4]:
```

	date	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	
0	2008-06-01 01:00:00	119.00	0.47	119.00	119.00	119.00	83.089996	120.699997	119.00	16.95
1	2008-06-01 01:00:00	119.00	0.59	119.00	119.00	119.00	94.820000	130.399994	119.00	17.46
2	2008-06-01 01:00:00	119.00	0.55	119.00	119.00	119.00	75.919998	104.599998	119.00	13.41
3	2008-06-01 01:00:00	119.00	0.36	119.00	119.00	119.00	61.029999	66.559998	119.00	23.11
4	2008-06-01 01:00:00	1.68	0.80	1.70	3.01	0.30	105.199997	214.899994	1.61	12.11
...
226387	2008-11-01 00:00:00	0.48	0.30	0.57	1.00	0.31	13.050000	14.160000	0.91	57.46
226388	2008-11-01 00:00:00	119.00	0.30	119.00	119.00	119.00	41.880001	48.500000	119.00	35.81
226389	2008-11-01 00:00:00	0.25	119.00	0.56	119.00	0.11	83.610001	102.199997	119.00	14.11
226390	2008-11-01 00:00:00	0.54	119.00	2.70	119.00	0.18	70.639999	81.860001	119.00	119.00
226391	2008-11-01 00:00:00	0.75	0.36	1.20	2.75	0.16	58.240002	74.239998	1.64	31.91

226392 rows × 17 columns

```
In [5]:
```

```
Out[5]: Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',  
'PM10', 'PM25', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'],  
             dtype='object')
```

```
In [6]: c=b.head(11)
```

```
Out[6]:
```

	date	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3
0	2008-06-01 01:00:00	119.00	0.47	119.00	119.00	119.00	83.089996	120.699997	119.00	16.990000
1	2008-06-01 01:00:00	119.00	0.59	119.00	119.00	119.00	94.820000	130.399994	119.00	17.469999
2	2008-06-01 01:00:00	119.00	0.55	119.00	119.00	119.00	75.919998	104.599998	119.00	13.470000
3	2008-06-01 01:00:00	119.00	0.36	119.00	119.00	119.00	61.029999	66.559998	119.00	23.110001
4	2008-06-01 01:00:00	1.68	0.80	1.70	3.01	0.30	105.199997	214.899994	1.61	12.120000
5	2008-06-01 01:00:00	119.00	0.47	119.00	119.00	0.22	67.820000	101.099998	119.00	20.610001
6	2008-06-01 01:00:00	0.17	0.40	0.44	119.00	0.15	72.639999	91.220001	119.00	17.040001
7	2008-06-01 01:00:00	119.00	0.51	119.00	119.00	119.00	80.440002	141.500000	119.00	10.310000
8	2008-06-01 01:00:00	119.00	0.36	119.00	119.00	119.00	68.150002	85.639999	119.00	23.580000
9	2008-06-01 01:00:00	119.00	0.18	119.00	119.00	0.16	58.330002	64.769997	119.00	35.060001
10	2008-06-01 01:00:00	119.00	0.45	119.00	119.00	119.00	53.700001	66.610001	119.00	27.180000


```
In [7]: d=c[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3']]
```

```
Out[7]:
```

	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3
0	119.00	0.47	119.00	119.00	119.00	83.089996	120.699997	119.00	16.990000
1	119.00	0.59	119.00	119.00	119.00	94.820000	130.399994	119.00	17.469999
2	119.00	0.55	119.00	119.00	119.00	75.919998	104.599998	119.00	13.470000
3	119.00	0.36	119.00	119.00	119.00	61.029999	66.559998	119.00	23.110001
4	1.68	0.80	1.70	3.01	0.30	105.199997	214.899994	1.61	12.120000
5	119.00	0.47	119.00	119.00	0.22	67.820000	101.099998	119.00	20.610001
6	0.17	0.40	0.44	119.00	0.15	72.639999	91.220001	119.00	17.040001
7	119.00	0.51	119.00	119.00	119.00	80.440002	141.500000	119.00	10.310000
8	119.00	0.36	119.00	119.00	119.00	68.150002	85.639999	119.00	23.580000
9	119.00	0.18	119.00	119.00	0.16	58.330002	64.769997	119.00	35.060001
10	119.00	0.45	119.00	119.00	119.00	53.700001	66.610001	119.00	27.180000


In [8]:

Out[8]: <AxesSubplot:>



In [9]:

Out[9]: <seaborn.axisgrid.PairGrid at 0x11e6de54280>

In [10]: `x=d[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY']]`In [11]: `from sklearn.model_selection import train_test_split`In [12]: `from sklearn.linear_model import LinearRegression`
`lr=LinearRegression()`Out[12]: `LinearRegression()`

In [13]:

-7.105427357601002e-14


```
In [14]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
```

```
Out[14]:
```

	Co-efficient
BEN	2.989965e-16
CO	9.585229e-14
EBE	-6.116106e-16
MXY	-5.266844e-16
NMHC	2.110477e-15
NO_2	7.048595e-16
NOx	1.000000e+00
OXY	-5.520637e-16

```
In [15]: prediction=lr.predict(x_test)
```

```
Out[15]: <matplotlib.collections.PathCollection at 0x11e723ed4c0>
```



```
In [16]:
```

```
1.0
```

```
In [17]:
```

```
In [18]: rr=Ridge(alpha=10)
```

```
Out[18]: Ridge(alpha=10)
```

```
In [19]:
```

```
Out[19]: 0.99994569566436
```

```
In [20]: la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

```
Out[20]: Lasso(alpha=10)
```

In [21]:

Out[21]: 0.9999665160868221

In [22]: a1=b.head(6000)

Out[22]:

	date	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3
0	2008-06-01 01:00:00	119.00	0.47	119.0	119.00	119.0	83.089996	120.699997	119.00	16.990000
1	2008-06-01 01:00:00	119.00	0.59	119.0	119.00	119.0	94.820000	130.399994	119.00	17.469999
2	2008-06-01 01:00:00	119.00	0.55	119.0	119.00	119.0	75.919998	104.599998	119.00	13.470000
3	2008-06-01 01:00:00	119.00	0.36	119.0	119.00	119.0	61.029999	66.559998	119.00	23.110000
4	2008-06-01 01:00:00	1.68	0.80	1.7	3.01	0.3	105.199997	214.899994	1.61	12.120000
...
5995	2008-06-10 15:00:00	119.00	0.27	119.0	119.00	119.0	32.000000	35.639999	119.00	67.980000
5996	2008-06-10 15:00:00	119.00	0.32	119.0	119.00	119.0	45.299999	57.360001	119.00	59.110000
5997	2008-06-10 15:00:00	119.00	0.24	119.0	119.00	119.0	26.160000	39.930000	119.00	52.200000
5998	2008-06-10 15:00:00	119.00	119.00	119.0	119.00	119.0	119.000000	119.000000	119.00	119.000000
5999	2008-06-10 15:00:00	119.00	0.31	119.0	119.00	119.0	68.599998	94.559998	119.00	42.720000

6000 rows × 11 columns

In [23]: e=a1[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',

In [24]: f=e.iloc[:,0:14]

In [25]:

In [26]: logr=LogisticRegression(max_iter=10000)

Out[26]: LogisticRegression(max_iter=10000)

In [27]: from sklearn.model_selection import train_test_split

In [28]:


```
In [29]: prediction=logr.predict(i)
```

```
[28079003]
```

```
In [30]:
```

```
Out[30]: array([28079001, 28079003, 28079004, 28079006, 28079007, 28079008,
                28079009, 28079011, 28079012, 28079014, 28079015, 28079016,
                28079018, 28079019, 28079021, 28079022, 28079023, 28079024,
                28079025, 28079026, 28079027, 28079036, 28079038, 28079039,
                28079040, 28079099], dtype=int64)
```

```
In [31]:
```

```
Out[31]: 4.63437247861083e-27
```

```
In [32]:
```

```
Out[32]: 0.7955388500721003
```

```
In [33]:
```

```
Out[33]: 0.575
```

```
In [34]: from sklearn.linear_model import ElasticNet
          en=ElasticNet()
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_d
escent.py:530: ConvergenceWarning: Objective did not converge. You might want
to increase the number of iterations. Duality gap: 8.3059915445693, toleranc
e: 1.6844122929515772
      model = cd_fast.enet_coordinate_descent(
```

```
Out[34]: ElasticNet()
```

```
In [35]:
```

```
[-0.0802347  0.          -0.12234881 -0.02045458 -0.          0.
  0.99934409  0.22163918]
```

```
In [36]:
```

```
0.22489545373029785
```

```
In [37]: prediction=en.predict(x_test)
```

```
0.6320535313373075
```

```
In [38]: from sklearn.ensemble import RandomForestClassifier
          rfc=RandomForestClassifier()
```

```
Out[38]: RandomForestClassifier()
```

```
In [39]: parameters={'max_depth':[1,2,3,4,5],
                    'min_samples_leaf':[5,10,15,20,25],
                    'n_estimators':[10,20,30,40,50]}
```

```
In [40]: from sklearn.model_selection import GridSearchCV
         grid_search=GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="acc
```

```
Out[40]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
                    param_grid={'max_depth': [1, 2, 3, 4, 5],
                                'min_samples_leaf': [5, 10, 15, 20, 25],
                                'n_estimators': [10, 20, 30, 40, 50]},
                    scoring='accuracy')
```

```
In [41]:
```

```
Out[41]: 0.6535714285714286
```

```
In [42]:
```

```
In [43]: from sklearn.tree import plot_tree
         plt.figure(figsize=(80,50))
```

```
Out[43]: [Text(2178.346153846154, 2491.5, 'X[13] <= -0.225\ngini = 0.961\nsamples = 26
23\nvalue = [160, 159, 143, 166, 159, 171, 177, 179, 185, 125\n158, 154, 152,
166, 148, 185, 153, 188, 160, 151\n179, 143, 153, 163, 163, 160]'),
         Text(1244.7692307692307, 2038.5, 'X[11] <= -0.06\ngini = 0.874\nsamples = 81
9\nvalue = [0, 0, 0, 164, 0, 169, 0, 0, 0, 0, 157, 0, 0\n0, 0, 153, 188,
0, 151, 179, 0, 0, 0, 0\n160]'),
         Text(686.7692307692307, 1585.5, 'X[5] <= -0.845\ngini = 0.763\nsamples = 37
4\nvalue = [0, 0, 0, 27, 0, 143, 0, 0, 0, 0, 1, 0, 0, 0\n0, 0, 143, 176, 0,
0, 0, 0, 0, 0, 0, 101]'),
         Text(343.38461538461536, 1132.5, 'X[2] <= -1.505\ngini = 0.44\nsamples = 13
0\nvalue = [0, 0, 0, 15, 0, 5, 0, 0, 0, 0, 0, 1, 0, 0, 0\n0, 0, 10, 149, 0, 0,
0, 0, 0, 0, 0, 23]'),
         Text(171.69230769230768, 679.5, 'X[8] <= 0.585\ngini = 0.727\nsamples = 52\n
value = [0, 0, 0, 12, 0, 5, 0, 0, 0, 0, 0, 0, 0, 0, 0\n0, 0, 8, 30, 0, 0, 0, 0,
0, 0, 0, 23]'),
         Text(85.84615384615384, 226.5, 'gini = 0.728\nsamples = 35\nvalue = [0, 0,
0, 12, 0, 5, 0, 0, 0, 0, 0, 0, 0, 0, 0\n0, 0, 5, 10, 0, 0, 0, 0, 0, 0, 0, 23]'),
         Text(257.53846153846155, 226.5, 'gini = 0.227\nsamples = 17\nvalue = [0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0\n0, 0, 3, 20, 0, 0, 0, 0, 0, 0, 0, 0]'),
         Text(515.0769230769231, 679.5, 'X[11] <= -0.06\ngini = 0.874\nsamples = 81
9\nvalue = [0, 0, 0, 164, 0, 169, 0, 0, 0, 0, 157, 0, 0\n0, 0, 153, 188,
0, 151, 179, 0, 0, 0, 0\n160]')]
```

From this observation I had observe that the LASSO is a highest accuracy of 0.9999665160868221

```
In [ ]:
```

