

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
import re
from sklearn.datasets import load_digits
```

```
In [2]: a=pd.read_csv(r"C:\Users\user\Downloads\C10_air\csvs_per_year\csvs_per_year\ma
```

Out[2]:

	date	BEN	CO	EBE	MXV	NMHC	NO_2	NOx	OXY	O_3	
0	2009-10-01 01:00:00	NaN	0.27	NaN	NaN	NaN	39.889999	48.150002	NaN	50.680000	18.2
1	2009-10-01 01:00:00	NaN	0.22	NaN	NaN	NaN	21.230000	24.260000	NaN	55.880001	10.5
2	2009-10-01 01:00:00	NaN	0.18	NaN	NaN	NaN	31.230000	34.880001	NaN	49.060001	25.1
3	2009-10-01 01:00:00	0.95	0.33	1.43	2.68	0.25	55.180000	81.360001	1.57	36.669998	26.5
4	2009-10-01 01:00:00	NaN	0.41	NaN	NaN	0.12	61.349998	76.260002	NaN	38.090000	23.7
...
215683	2009-06-01 00:00:00	0.50	0.22	0.39	0.75	0.09	22.000000	24.510000	1.00	82.239998	10.8
215684	2009-06-01 00:00:00	NaN	0.31	NaN	NaN	NaN	76.110001	101.099998	NaN	41.220001	9.9
215685	2009-06-01 00:00:00	0.13	NaN	0.86	NaN	0.23	81.050003	99.849998	NaN	24.830000	12.4
215686	2009-06-01 00:00:00	0.21	NaN	2.96	NaN	0.10	72.419998	82.959999	NaN	NaN	13.0
215687	2009-06-01 00:00:00	0.37	0.32	0.99	1.36	0.14	54.290001	64.480003	1.06	56.919998	15.3

215688 rows × 17 columns

In [3]:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 215688 entries, 0 to 215687
Data columns (total 17 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        215688 non-null  object
1   BEN         60082 non-null   float64
2   CO          190801 non-null  float64
3   EBE         60081 non-null   float64
4   MXY         24846 non-null   float64
5   NMHC        74748 non-null   float64
6   NO_2        214562 non-null  float64
7   NOx         214565 non-null  float64
8   OXY         24854 non-null   float64
9   O_3         204482 non-null  float64
10  PM10        196331 non-null  float64
11  PM25        55822 non-null   float64
12  PXY         24854 non-null   float64
13  SO_2        212671 non-null  float64
14  TCH         75213 non-null   float64
15  TOL         59920 non-null   float64
16  station     215688 non-null  int64
dtypes: float64(15), int64(1), object(1)
memory usage: 28.0+ MB
```

In [4]: `b=a.fillna(value=98)`

Out[4]:

	date	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3
0	2009-10-01 01:00:00	98.00	0.27	98.00	98.00	98.00	39.889999	48.150002	98.00	50.680000
1	2009-10-01 01:00:00	98.00	0.22	98.00	98.00	98.00	21.230000	24.260000	98.00	55.880001
2	2009-10-01 01:00:00	98.00	0.18	98.00	98.00	98.00	31.230000	34.880001	98.00	49.060001
3	2009-10-01 01:00:00	0.95	0.33	1.43	2.68	0.25	55.180000	81.360001	1.57	36.669998
4	2009-10-01 01:00:00	98.00	0.41	98.00	98.00	0.12	61.349998	76.260002	98.00	38.090000
...
215683	2009-06-01 00:00:00	0.50	0.22	0.39	0.75	0.09	22.000000	24.510000	1.00	82.239998
215684	2009-06-01 00:00:00	98.00	0.31	98.00	98.00	98.00	76.110001	101.099998	98.00	41.220001
215685	2009-06-01 00:00:00	0.13	98.00	0.86	98.00	0.23	81.050003	99.849998	98.00	24.830000
215686	2009-06-01 00:00:00	0.21	98.00	2.96	98.00	0.10	72.419998	82.959999	98.00	98.000000
215687	2009-06-01 00:00:00	0.37	0.32	0.99	1.36	0.14	54.290001	64.480003	1.06	56.919998

215688 rows × 17 columns

In [5]:

Out[5]: `Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
'PM10', 'PM25', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'],
 dtype='object')`

```
In [6]: c=b.head(11)
```

Out[6]:

	date	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3	PM
0	2009-10-01 01:00:00	98.00	0.27	98.00	98.00	98.00	39.889999	48.150002	98.00	50.680000	18.2600
1	2009-10-01 01:00:00	98.00	0.22	98.00	98.00	98.00	21.230000	24.260000	98.00	55.880001	10.5800
2	2009-10-01 01:00:00	98.00	0.18	98.00	98.00	98.00	31.230000	34.880001	98.00	49.060001	25.1900
3	2009-10-01 01:00:00	0.95	0.33	1.43	2.68	0.25	55.180000	81.360001	1.57	36.669998	26.5300
4	2009-10-01 01:00:00	98.00	0.41	98.00	98.00	0.12	61.349998	76.260002	98.00	38.090000	23.7600
5	2009-10-01 01:00:00	98.00	0.29	98.00	98.00	98.00	43.200001	50.080002	98.00	35.840000	21.8700
6	2009-10-01 01:00:00	98.00	0.20	98.00	98.00	98.00	35.430000	38.520000	98.00	33.549999	17.3500
7	2009-10-01 01:00:00	98.00	0.15	98.00	98.00	98.00	27.309999	33.150002	98.00	53.549999	16.5200
8	2009-10-01 01:00:00	98.00	0.21	98.00	98.00	0.39	33.889999	40.799999	98.00	58.549999	16.6500
9	2009-10-01 01:00:00	98.00	0.32	98.00	98.00	98.00	46.349998	60.540001	98.00	45.340000	15.1600
10	2009-10-01 01:00:00	98.00	0.24	98.00	98.00	98.00	30.860001	35.590000	98.00	56.520000	14.4200

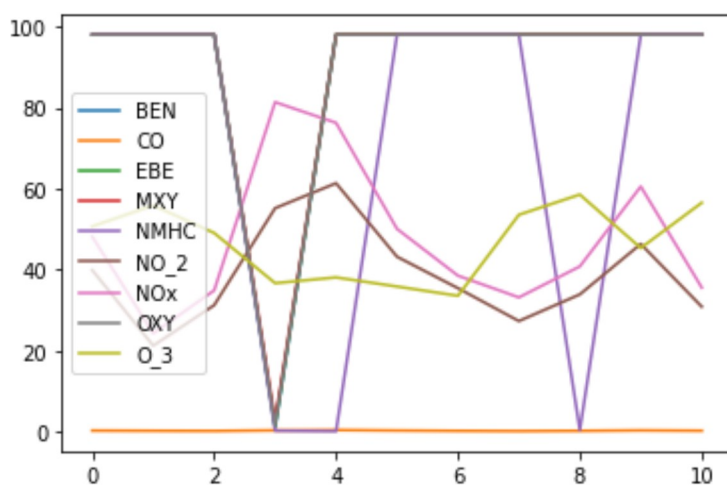
```
In [7]: d=c[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3']]
```

Out[7]:

	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3
0	98.00	0.27	98.00	98.00	98.00	39.889999	48.150002	98.00	50.680000
1	98.00	0.22	98.00	98.00	98.00	21.230000	24.260000	98.00	55.880001
2	98.00	0.18	98.00	98.00	98.00	31.230000	34.880001	98.00	49.060001
3	0.95	0.33	1.43	2.68	0.25	55.180000	81.360001	1.57	36.669998
4	98.00	0.41	98.00	98.00	0.12	61.349998	76.260002	98.00	38.090000
5	98.00	0.29	98.00	98.00	98.00	43.200001	50.080002	98.00	35.840000
6	98.00	0.20	98.00	98.00	98.00	35.430000	38.520000	98.00	33.549999
7	98.00	0.15	98.00	98.00	98.00	27.309999	33.150002	98.00	53.549999
8	98.00	0.21	98.00	98.00	0.39	33.889999	40.799999	98.00	58.549999
9	98.00	0.32	98.00	98.00	98.00	46.349998	60.540001	98.00	45.340000
10	98.00	0.24	98.00	98.00	98.00	30.860001	35.590000	98.00	56.520000


In [8]:

Out[8]: <AxesSubplot:>



In [9]:

Out[9]: <seaborn.axisgrid.PairGrid at 0x15228f57eb0>

In [10]: `x=d[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY']]`In [11]: `from sklearn.model_selection import train_test_split`In [12]: `from sklearn.linear_model import LinearRegression
lr=LinearRegression()`

Out[12]: LinearRegression()

In [13]:

-2.1316282072803006e-14


```
In [14]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
```

```
Out[14]:
```

	Co-efficient
BEN	1.002284e-15
CO	-4.236772e-14
EBE	-1.779152e-16
MXY	-6.386554e-16
NMHC	-6.769435e-17
NO_2	1.006278e-15
NOx	1.000000e+00
OXY	-5.039871e-17

```
In [15]: prediction=lr.predict(x_test)
```

```
Out[15]: <matplotlib.collections.PathCollection at 0x1522c531550>
```



```
In [16]:
```

```
1.0
```

```
In [17]:
```

```
In [18]: rr=Ridge(alpha=10)
```

```
Out[18]: Ridge(alpha=10)
```

```
In [19]:
```

```
Out[19]: 0.9932959728090598
```

```
In [20]: la=Lasso(alpha=10)
```

```
la.fit(x_train,y_train)
```

```
Out[20]: Lasso(alpha=10)
```

In [21]:

Out[21]: 0.9980831728379033

In [22]: a1=b.head(3000)

Out[22]:

	date	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3	
0	2009-10-01 01:00:00	98.00	0.27	98.00	98.00	98.00	39.889999	48.150002	98.00	50.680000	1
1	2009-10-01 01:00:00	98.00	0.22	98.00	98.00	98.00	21.230000	24.260000	98.00	55.880001	1
2	2009-10-01 01:00:00	98.00	0.18	98.00	98.00	98.00	31.230000	34.880001	98.00	49.060001	2
3	2009-10-01 01:00:00	0.95	0.33	1.43	2.68	0.25	55.180000	81.360001	1.57	36.669998	2
4	2009-10-01 01:00:00	98.00	0.41	98.00	98.00	0.12	61.349998	76.260002	98.00	38.090000	2
...
2995	2009-10-06 00:00:00	1.10	0.75	1.16	2.85	0.66	128.300003	192.100006	0.92	0.600000	5
2996	2009-10-06 00:00:00	98.00	0.33	98.00	98.00	98.00	88.220001	108.199997	98.00	20.520000	4
2997	2009-10-06 00:00:00	1.10	98.00	0.39	98.00	0.42	124.400002	221.800003	98.00	6.190000	5
2998	2009-10-06 00:00:00	2.75	98.00	1.55	98.00	0.25	129.399994	166.000000	98.00	98.000000	2
2999	2009-10-06 00:00:00	2.23	0.83	2.36	7.03	0.44	126.300003	246.699997	3.18	8.110000	6

3000 rows × 17 columns

In [23]: e=a1[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',

In [24]: f=e.iloc[:,0:14]

In [25]:

In [26]: logr=LogisticRegression(max_iter=10000)

Out[26]: LogisticRegression(max_iter=10000)

In [27]: from sklearn.model_selection import train_test_split

In [28]:


```
In [29]: prediction=logr.predict(i)
```

```
[28079021]
```

```
In [30]:
```

```
Out[30]: array([28079003, 28079004, 28079006, 28079007, 28079008, 28079009,
                28079011, 28079012, 28079014, 28079016, 28079017, 28079018,
                28079019, 28079021, 28079022, 28079023, 28079024, 28079025,
                28079026, 28079027, 28079036, 28079038, 28079039, 28079040,
                28079099], dtype=int64)
```

```
In [31]:
```

```
Out[31]: 5.665341540178148e-99
```

```
In [32]:
```

```
Out[32]: 1.705432077416329e-33
```

```
In [33]:
```

```
Out[33]: 0.5888888888888889
```

```
In [34]: from sklearn.linear_model import ElasticNet
          en=ElasticNet()
```

```
Out[34]: ElasticNet()
```

```
In [35]:
```

```
[ -0.          0.          -0.01381274 -0.0013857  -0.00329021  0.13429634
  0.88971371 -0.          ]
```

```
In [36]:
```

```
1.5507170991143298
```

```
In [37]: prediction=en.predict(x_test)
```

```
0.9988450094456112
```

```
In [38]: from sklearn.ensemble import RandomForestClassifier
          rfc=RandomForestClassifier()
```

```
Out[38]: RandomForestClassifier()
```

```
In [39]: parameters={'max_depth':[1,2,3,4,5],
                    'min_samples_leaf':[5,10,15,20,25],
                    'n_estimators':[10,20,30,40,50]}
```

```
Out[40]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
                      param_grid={'max_depth': [1, 2, 3, 4, 5],
                                   'min_samples_leaf': [5, 10, 15, 20, 25],
                                   'n_estimators': [10, 20, 30, 40, 50]},
                      scoring='accuracy')
```

Out[41]: 0.5780952380952381

```
Out[43]: [Text(1923.3191489361702, 2491.5, 'X[10] <= -0.853\ngini = 0.96\nsamples = 13
21\nvalue = [87, 94, 80, 90, 92, 88, 81, 85, 71, 89, 72, 81\n90, 104, 87, 73,
65, 87, 88, 77, 76, 79, 90, 83\n91]'),
Text(831.063829787234, 2038.5, 'X[7] <= -2.742\ngini = 0.661\nsamples = 152\n
value = [0, 0, 80, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0\n0, 0, 65, 0, 0, 0, 0, 0,
0, 0, 91]'),
Text(379.9148936170213, 1585.5, 'X[5] <= -0.393\ngini = 0.316\nsamples = 52\n
value = [0, 0, 7, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0\n0, 0, 58, 0, 0, 0, 0, 0,
0, 0, 6]'),
Text(189.95744680851064, 1132.5, 'X[2] <= -1.637\ngini = 0.166\nsamples = 3
9\nvalue = [0, 0, 3, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0\n0, 0, 51, 0, 0, 0, 0,
0, 0, 0, 2]'),
Text(94.97872340425532, 679.5, 'gini = 0.0\nsamples = 27\nvalue = [0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0\n0, 0, 38, 0, 0, 0, 0, 0, 0, 0, 0]'),
Text(284.93617021276594, 679.5, 'X[6] <= -0.816\ngini = 0.438\nsamples = 12\n
value = [0, 0, 3, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0\n0, 0, 13, 0, 0, 0, 0, 0,
0, 0, 2]'),
Text(189.95744680851064, 226.5, 'gini = 0.625\nsamples = 7\nvalue = [0, 0,
2, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0\n0, 0, 4, 0, 0, 0, 0, 0, 0, 0, 2]'),
Text(270.2148936170213, 226.5, 'gini = 0.18\nsamples = 5\nvalue = [0, 0, 1,
```

From this observation I had observe that the ELASTICNET is a highest accuracy of 0.9988450094456112

In []: