

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
import re
from sklearn.datasets import load_digits
```

```
In [2]: a=pd.read_csv(r"C:\Users\user\Downloads\C10_air\csvs_per_year\csvs_per_year\ma
```

Out[2]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL
0	2012-09-01 01:00:00	NaN	0.2	NaN	NaN	7.0	18.0	NaN	NaN	NaN	2.0	NaN	NaN
1	2012-09-01 01:00:00	0.3	0.3	0.7	NaN	3.0	18.0	55.0	10.0	9.0	1.0	NaN	2.4
2	2012-09-01 01:00:00	0.4	NaN	0.7	NaN	2.0	10.0	NaN	NaN	NaN	NaN	NaN	1.5
3	2012-09-01 01:00:00	NaN	0.2	NaN	NaN	1.0	6.0	50.0	NaN	NaN	NaN	NaN	NaN
4	2012-09-01 01:00:00	NaN	NaN	NaN	NaN	1.0	13.0	54.0	NaN	NaN	3.0	NaN	NaN
...
210715	2012-03-01 00:00:00	NaN	0.6	NaN	NaN	37.0	84.0	14.0	NaN	NaN	NaN	NaN	NaN
210716	2012-03-01 00:00:00	NaN	0.4	NaN	NaN	5.0	76.0	NaN	17.0	NaN	7.0	NaN	NaN
210717	2012-03-01 00:00:00	NaN	NaN	NaN	0.34	3.0	41.0	24.0	NaN	NaN	NaN	1.34	NaN
210718	2012-03-01 00:00:00	NaN	NaN	NaN	NaN	2.0	44.0	36.0	NaN	NaN	NaN	NaN	NaN
210719	2012-03-01 00:00:00	NaN	NaN	NaN	NaN	2.0	56.0	40.0	18.0	NaN	NaN	NaN	NaN

210720 rows × 14 columns

In [3]:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 210720 entries, 0 to 210719
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        210720 non-null object
1   BEN         51511 non-null  float64
2   CO          87097 non-null  float64
3   EBE         51482 non-null  float64
4   NMHC        30736 non-null  float64
5   NO          209871 non-null float64
6   NO_2        209872 non-null float64
7   O_3         122339 non-null float64
8   PM10        104838 non-null float64
9   PM25        52164 non-null  float64
10  SO_2        87333 non-null  float64
11  TCH         30736 non-null  float64
12  TOL         51373 non-null  float64
13  station     210720 non-null int64
dtypes: float64(12), int64(1), object(1)
memory usage: 22.5+ MB
```

In [4]: `b=a.fillna(value=188)`

Out[4]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH
0	2012-09-01 01:00:00	188.0	0.2	188.0	188.00	7.0	18.0	188.0	188.0	188.0	2.0	188.00
1	2012-09-01 01:00:00	0.3	0.3	0.7	188.00	3.0	18.0	55.0	10.0	9.0	1.0	188.00
2	2012-09-01 01:00:00	0.4	188.0	0.7	188.00	2.0	10.0	188.0	188.0	188.0	188.0	188.00
3	2012-09-01 01:00:00	188.0	0.2	188.0	188.00	1.0	6.0	50.0	188.0	188.0	188.0	188.00
4	2012-09-01 01:00:00	188.0	188.0	188.0	188.00	1.0	13.0	54.0	188.0	188.0	3.0	188.00
...
210715	2012-03-01 00:00:00	188.0	0.6	188.0	188.00	37.0	84.0	14.0	188.0	188.0	188.0	188.00
210716	2012-03-01 00:00:00	188.0	0.4	188.0	188.00	5.0	76.0	188.0	17.0	188.0	7.0	188.00
210717	2012-03-01 00:00:00	188.0	188.0	188.0	0.34	3.0	41.0	24.0	188.0	188.0	188.0	1.34
210718	2012-03-01 00:00:00	188.0	188.0	188.0	188.00	2.0	44.0	36.0	188.0	188.0	188.0	188.00
210719	2012-03-01 00:00:00	188.0	188.0	188.0	188.00	2.0	56.0	40.0	18.0	188.0	188.0	188.00

210720 rows × 14 columns

In [5]:

Out[5]: `Index(['date', 'BEN', 'CO', 'EBE', 'NMHC', 'NO', 'NO_2', 'O_3', 'PM10', 'PM25', 'SO_2', 'TCH', 'TOL', 'station'], dtype='object')`

```
In [6]: c=b.head(11)
```

```
Out[6]:
```

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TOL
0	2012-09-01 01:00:00	188.0	0.2	188.0	188.00	7.0	18.0	188.0	188.0	188.0	2.0	188.00	188.0
1	2012-09-01 01:00:00	0.3	0.3	0.7	188.00	3.0	18.0	55.0	10.0	9.0	1.0	188.00	2.4
2	2012-09-01 01:00:00	0.4	188.0	0.7	188.00	2.0	10.0	188.0	188.0	188.0	188.0	188.00	1.5
3	2012-09-01 01:00:00	188.0	0.2	188.0	188.00	1.0	6.0	50.0	188.0	188.0	188.0	188.00	188.0
4	2012-09-01 01:00:00	188.0	188.0	188.0	188.00	1.0	13.0	54.0	188.0	188.0	3.0	188.00	188.0
5	2012-09-01 01:00:00	0.2	0.2	1.0	188.00	1.0	9.0	57.0	14.0	188.0	1.0	188.00	0.2
6	2012-09-01 01:00:00	0.4	0.2	0.8	0.24	1.0	7.0	57.0	11.0	7.0	2.0	1.33	0.6
7	2012-09-01 01:00:00	188.0	188.0	188.0	0.11	1.0	2.0	65.0	188.0	188.0	188.0	1.18	188.0
8	2012-09-01 01:00:00	188.0	0.2	188.0	188.00	6.0	14.0	57.0	188.0	188.0	2.0	188.00	188.0
9	2012-09-01 01:00:00	188.0	0.2	188.0	188.00	1.0	7.0	188.0	13.0	188.0	1.0	188.00	188.0
10	2012-09-01 01:00:00	0.2	188.0	0.7	188.00	3.0	13.0	188.0	12.0	6.0	1.0	188.00	0.6


```
In [7]: d=c[['BEN', 'CO', 'EBE', 'NMHC', 'NO_2', 'O_3']]
```

```
Out[7]:
```

	BEN	CO	EBE	NMHC	NO_2	O_3
0	188.0	0.2	188.0	188.00	18.0	188.0
1	0.3	0.3	0.7	188.00	18.0	55.0
2	0.4	188.0	0.7	188.00	10.0	188.0
3	188.0	0.2	188.0	188.00	6.0	50.0
4	188.0	188.0	188.0	188.00	13.0	54.0
5	0.2	0.2	1.0	188.00	9.0	57.0
6	0.4	0.2	0.8	0.24	7.0	57.0
7	188.0	188.0	188.0	0.11	2.0	65.0
8	188.0	0.2	188.0	188.00	14.0	57.0
9	188.0	0.2	188.0	188.00	7.0	188.0
10	0.2	188.0	0.7	188.00	13.0	188.0


In [8]:

Out[8]: <AxesSubplot:>



In [9]:

Out[9]: <seaborn.axisgrid.PairGrid at 0x1b8df750610>

In [10]: `x=d[['BEN', 'CO', 'EBE', 'NMHC', 'NO_2']]`In [11]: `from sklearn.model_selection import train_test_split`In [12]: `from sklearn.linear_model import LinearRegression`
`lr=LinearRegression()`

Out[12]: LinearRegression()

In [13]:

2.2431834167946363e-11


```
In [14]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
```

```
Out[14]:
```

	Co-efficient
BEN	5.588715e-11
CO	1.000000e+00
EBE	-5.600700e-11
NMHC	4.406879e-16
NO_2	-2.105775e-15

```
In [15]: prediction=lr.predict(x_test)
```

```
Out[15]: <matplotlib.collections.PathCollection at 0x1b8e27d0310>
```



```
In [16]:
```

```
1.0
```

```
In [17]:
```

```
In [18]: rr=Ridge(alpha=10)
```

```
Out[18]: Ridge(alpha=10)
```

```
In [19]:
```

```
Out[19]: 0.9999995643541453
```

```
In [20]: la=Lasso(alpha=10)
```

```
Out[20]: Lasso(alpha=10)
```

```
In [21]:
```

```
Out[21]: 0.9999977145129909
```

In [22]: `a1=b.head(5000)`

Out[22]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	T
0	2012-09-01 01:00:00	188.0	0.2	188.0	188.00	7.0	18.0	188.0	188.0	188.0	2.0	188.00	18
1	2012-09-01 01:00:00	0.3	0.3	0.7	188.00	3.0	18.0	55.0	10.0	9.0	1.0	188.00	:
2	2012-09-01 01:00:00	0.4	188.0	0.7	188.00	2.0	10.0	188.0	188.0	188.0	188.0	188.00	
3	2012-09-01 01:00:00	188.0	0.2	188.0	188.00	1.0	6.0	50.0	188.0	188.0	188.0	188.00	18
4	2012-09-01 01:00:00	188.0	188.0	188.0	188.00	1.0	13.0	54.0	188.0	188.0	3.0	188.00	18
...	
4995	2012-09-09 17:00:00	188.0	0.2	188.0	188.00	2.0	8.0	96.0	188.0	188.0	188.0	188.00	18
4996	2012-09-09 17:00:00	188.0	188.0	188.0	188.00	2.0	5.0	99.0	188.0	188.0	3.0	188.00	18
4997	2012-09-09 17:00:00	0.2	0.2	1.0	188.00	1.0	5.0	93.0	27.0	188.0	1.0	188.00	:
4998	2012-09-09 17:00:00	0.5	0.2	1.1	0.22	1.0	3.0	97.0	27.0	12.0	3.0	1.32	:
4999	2012-09-09 17:00:00	188.0	188.0	188.0	0.11	1.0	4.0	110.0	188.0	188.0	188.0	1.18	18

5000 rows × 14 columns

In [23]: `e=a1[['BEN', 'CO', 'EBE', 'NMHC', 'NO_2', 'O_3',`

In [24]: `f=e.iloc[:,0:14]`

In [25]: `from sklearn.linear_model import LogisticRegression`

In [26]: `logr=LogisticRegression(max_iter=10000)`

Out[26]: `LogisticRegression(max_iter=10000)`

In [27]: `from sklearn.model_selection import train_test_split`

In [28]: `X_train, X_test, y_train, y_test = train_test_split(f, a1['T'],`

In [29]: `prediction=logr.predict(i)`

`[28079050]`

In [30]:

```
Out[30]: array([28079004, 28079008, 28079011, 28079016, 28079017, 28079018,
                28079024, 28079027, 28079035, 28079036, 28079038, 28079039,
                28079040, 28079047, 28079048, 28079049, 28079050, 28079054,
                28079055, 28079056, 28079057, 28079058, 28079059, 28079060],
                dtype=int64)
```

In [31]:

Out[31]: 0.0

In [32]:

Out[32]: 0.0

In [33]:

Out[33]: 0.95733333333333334

```
In [34]: from sklearn.linear_model import ElasticNet
          en=ElasticNet()
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:530: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 5.37854093374016, tolerance: 5.037333428315741
  model = cd_fast.enet_coordinate_descent(
```

Out[34]: ElasticNet()

In [35]:

```
[ 2.18053242e-01  9.99868788e-01 -2.18377624e-01 -1.22495536e-05
 0.00000000e+00]
```

In [36]:

0.07726383756725852

```
In [37]: prediction=en.predict(x_test)
```

0.9999996066476035

```
In [38]: from sklearn.ensemble import RandomForestClassifier
         rfc=RandomForestClassifier()
```

```
Out[38]: RandomForestClassifier()
```

```
In [39]: parameters={'max_depth':[1,2,3,4,5],
                    'min_samples_leaf':[5,10,15,20,25],
                    'n_estimators':[10,20,30,40,50]}
```

```
In [40]: from sklearn.model_selection import GridSearchCV
         grid_search=GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="acc
```

```
Out[40]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
                    param_grid={'max_depth': [1, 2, 3, 4, 5],
                                'min_samples_leaf': [5, 10, 15, 20, 25],
                                'n_estimators': [10, 20, 30, 40, 50]},
                    scoring='accuracy')
```

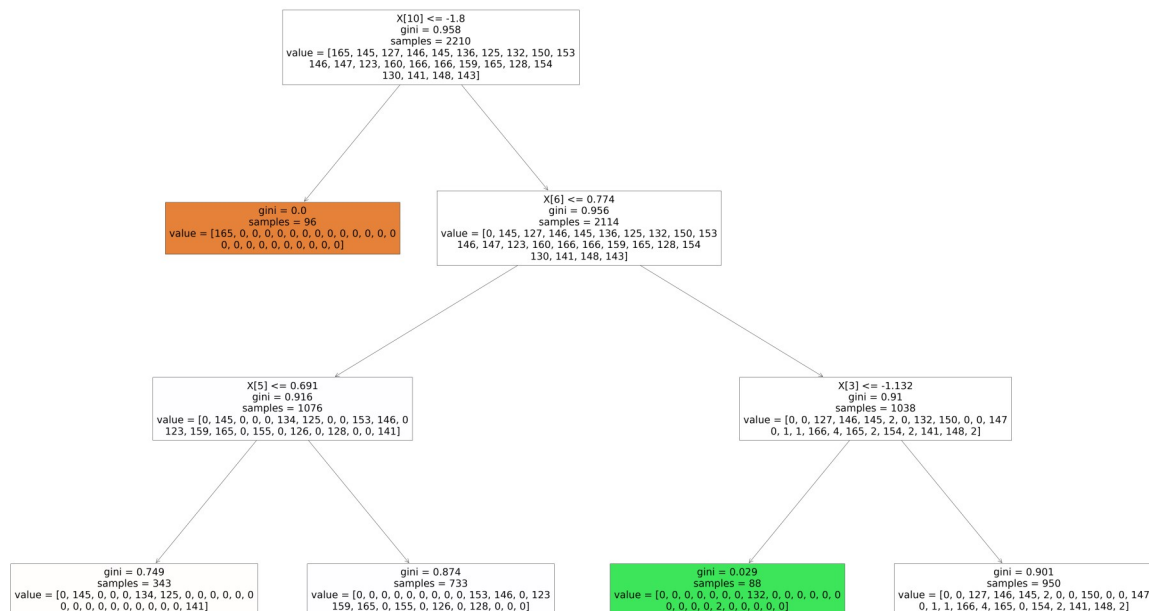
```
In [41]:
```

```
Out[41]: 0.9982857142857142
```

```
In [42]:
```

```
In [43]: from sklearn.tree import plot_tree
plt.figure(figsize=(80,50))
```

```
Out[43]: [Text(1674.0, 2378.25, 'X[10] <= -1.8\ngini = 0.958\nnsamples = 2210\nvalue = [165, 145, 127, 146, 145, 136, 125, 132, 150, 153\n146, 147, 123, 160, 166, 166, 159, 165, 128, 154\n130, 141, 148, 143]'),
Text(1116.0, 1698.75, 'gini = 0.0\nnsamples = 96\nvalue = [165, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]'),
Text(2232.0, 1698.75, 'X[6] <= 0.774\ngini = 0.956\nnsamples = 2114\nvalue = [0, 145, 127, 146, 145, 136, 125, 132, 150, 153\n146, 147, 123, 160, 166, 166, 159, 165, 128, 154\n6, 159, 165, 128, 154\n130, 141, 148, 143]'),
Text(1116.0, 1019.25, 'X[5] <= 0.691\ngini = 0.916\nnsamples = 1076\nvalue = [0, 145, 0, 0, 0, 134, 125, 0, 0, 153, 146, 0\n123, 159, 165, 0, 155, 0, 126, 0, 128, 0, 0, 141]'),
Text(558.0, 339.75, 'gini = 0.749\nnsamples = 343\nvalue = [0, 145, 0, 0, 0, 134, 125, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 141]'),
Text(1674.0, 339.75, 'gini = 0.874\nnsamples = 733\nvalue = [0, 0, 0, 0, 0, 153, 146, 0, 123\n159, 165, 0, 155, 0, 126, 0, 128, 0, 0, 0]'),
Text(3348.0, 1019.25, 'X[3] <= -1.132\ngini = 0.91\nnsamples = 1038\nvalue = [0, 0, 127, 146, 145, 2, 0, 132, 150, 0, 0, 147\n0, 1, 1, 166, 4, 165, 2, 154, 2, 141, 148, 2]'),
Text(2790.0, 339.75, 'gini = 0.029\nnsamples = 88\nvalue = [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]'),
Text(3906.0, 339.75, 'gini = 0.901\nnsamples = 950\nvalue = [0, 0, 127, 146, 145, 2, 0, 0, 150, 0, 0, 147\n0, 1, 1, 166, 4, 165, 0, 154, 2, 141, 148, 2]')]
```



From this observation I had observe that the ELASTICNET is a highest accuracy of 0.9999996066476035

In []: